RESUME OS

Master Intelligence System

v1.0 — PickCV Platform

Input: Resume + Job Description / Link

Output: 7-10 Optimised, Role-Aligned Resume Variants

Compounding intelligence. Zero hallucination.

Kanna Uppalapati + Adithya Saladi

Bhoomer Digital Solutions Pvt. Ltd. | PickCV Platform | March 2026

SECTION 0: PURPOSE AND GOVERNING PHILOSOPHY

This document is the master specification for the PickCV Resume OS — the intelligence system that runs every time a candidate submits a resume and a job description or job link.

It is not a formatting guide. It is not a keyword checklist. It is a decision-grade career signal engineering system.

> **The Central Problem Resume OS Solves**
>
> Talent is not missing. It is invisible to the systems that decide who gets hired. The resume is the signal layer between a candidate's capability and a recruiter's decision. Most resumes fail not because the candidate lacks ability — but because the signal is misrepresented, misformatted, or mistargeted. Resume OS corrects all three.

0.1 What This System Does

When a candidate provides their master resume and a target job (JD text or job link), Resume OS:

- Classifies the role across Function, Level, and Environment
- Detects the target ATS and applies system-specific formatting rules
- Performs gap analysis between the candidate's signal and the JD requirements
- Runs 8 optimization engines on the content
- Generates 7 to 10 structurally distinct resume variants — identical in factual content, different in structure, order, tone, and emphasis
- Scores each variant on ATS match, recruiter scan, and signal density
- Produces a strategy log explaining every change made

0.2 What This System Does NOT Do

- It does not fabricate experience, metrics, or achievements
- It does not add skills the candidate has not demonstrated
- It does not produce a generic 'improved' resume — it produces role-specific signal alignment
- It does not run once — it learns from every application and compounds intelligence over time

> **The Governing Principle**
>
> There is no universal perfect resume. There is no ATS hack. There is no keyword cheat code. There is only alignment with a specific problem. Resume OS engineers that alignment — precisely, systematically, and verifiably.

SECTION 1: CORE ARCHITECTURE

Resume OS processes every input through six sequential intelligence layers. Each layer outputs structured signals that feed the next layer. No layer is skipped.

EXECUTION FLOW

> INPUT: Resume + JD / Job Link
--> LAYER 1 — Role Intelligence: Classify function, level, environment, KPIs
--> LAYER 2 — ATS Intelligence: Detect ATS, apply formatting and keyword rules
--> LAYER 3 — Recruiter Decision Model: Apply 6-second scan and risk model
--> LAYER 4 — Gap Analysis: JD requirements vs resume signals
--> LAYER 5 — Optimization Engine (8 sub-engines): Rewrite, restructure, amplify
> LAYER 6 — Output Engine: Generate 7-10 variants + scores + strategy log

Each layer is described in full in the sections that follow. Every rule in every layer is a hard constraint, not a suggestion. The system applies all rules in sequence and does not override earlier layer outputs without justification logged in the strategy file.

SECTION 2: LAYER 1 — ROLE INTELLIGENCE

Every role is classified before any resume work begins. Without role classification, keyword stuffing and generic formatting are the only tools available. Role classification enables the entire system.

2.1 The Role DNA Framework

Every JD produces a Role DNA output defined by four axes:

FUNCTION	What the role does: GTM (Sales, Marketing, CS) | Tech (Engineering, Product, Data) | Business (Ops, Strategy, HR) | Support (Finance, Legal, Admin)
LEVEL	What the role expects: L1-L2 Execution | L3 Ownership | L4 Strategic | L5+ Business Impact
ENVIRONMENT	Where the role lives: Startup (speed + ownership) | SME (versatility) | MNC (process + specialisation)
KPIs	How success is measured: extracted directly from JD language and inferred from role type at that level

> **Role DNA Output Format**
> Function + Level + Environment + KPIsExample: TA Leader + L4 + Startup + Metrics-driven hiring, offer acceptance rate, time-to-fill

2.2 Function Axis — 10 Consolidated Role Clusters

After applying a 50% skill overlap merge rule, all market roles consolidate into 10 clusters. Roles are merged where core skills — tools, deliverables, stakeholder relationships, bullet logic — overlap at more than 50% at the relevant experience band. Roles that diverge below 50% overlap retain independent treatment.

ID	Cluster Name	Merged Roles	Shared Core Skills
C1	Software Engineer / Full-Stack / Backend / Frontend	SWE, Full-Stack, Backend, Frontend, Mobile	DSA, system design, version control, CI/CD, cloud deployment
C2	Data Analyst + Business Analyst	DA, BA, Reporting Analyst, BI Analyst	SQL, Excel/Sheets, dashboards, stakeholder reporting, insight derivation
C3	Data Scientist + ML Engineer	DS, MLE, AI Engineer, Research Scientist	Python, model building, experimentation, deployment pipelines
C4	Product Manager + Product Analyst	PM, APM, Product Analyst, TPM	Roadmapping, metrics, cross-functional alignment, user research
C5	DevOps + Cloud Engineer + SRE	DevOps, Cloud Eng, SRE, Platform Eng	Kubernetes, Docker, Terraform, CI/CD, cloud platforms, IaC
C6	Sales + Business Development	SDR, BDR, AE, AM, KAM, Enterprise Sales	CRM, pipeline management, outreach, quota achievement, negotiation
C7	Marketing + Growth + Content	Performance Mktg, Growth, Content, PMM, Brand	Campaign execution, analytics, channel ownership, content strategy

ID	Cluster Name	Merged Roles	Shared Core Skills
C8	HR + Talent Acquisition + People Ops	Recruiter, HRBP, TA Lead, People Ops	Hiring systems, employee lifecycle, HRIS, stakeholder management
C9	Operations + Supply Chain + Consulting	Ops Analyst, Consultant, SCM, Project Mgr	Process improvement, cross-functional execution, cost reduction
C10	UX/UI Designer + Product Designer	UX Designer, UI Designer, Product Designer	Design thinking, Figma, prototyping, user research, handoff

2.3 Level Axis — Resume Bias by Level

Level determines what the resume must prove. This is the most important axis. A candidate targeting L3 with an L2-style resume fails regardless of keyword match.

Level	Band	What It Proves	Resume Bias — Every Bullet Must Reflect This
L1–L2	0–2 yrs	**Tasks completed**	Built X using Y that does Z. Evidence of use, not familiarity. Scope: personal or small team.
L3	3–6 yrs	**Ownership of outcomes**	Owned X. Led redesign of Y. Every bullet must show a decision made, not just work delivered.
L4	7–12 yrs	**Systems and strategy**	Designed the system. Defined the standard. Led the initiative. Org-level impact is mandatory.
L5+	12+ yrs	**Business and org impact**	Built the team. Set the vision. Drove the outcome. Individual contributions are secondary to system creation.

2.4 Environment Axis — Hiring Bias by Company Type

Type	What They Hire For	Signal Words to Amplify	Signal Words to Reduce
Startup (0–200)	Speed, ownership, ambiguity tolerance, builder mindset	Built from scratch, owned end-to-end, shipped in X weeks, operated without playbook	Governance, process adherence, worked within framework
SME (200–2,000)	Versatility, reliability, cross-functional capability	Wore multiple hats, supported X and Y simultaneously, delivered across teams	Hyper-specialisation, deep domain at expense of breadth
MNC (2,000+)	Process fluency, specialisation, stakeholder management, governance	Managed across 3 regions, compliant with X, led structured program, scaled within framework	Ambiguity, chaos, no process, solo execution

2.5 Emerging and AI-Adjacent Role Variants

The 2025-2026 job market has produced a set of role variants that require adjusted treatment within their parent clusters. These roles carry additional signal requirements beyond the base cluster template.

AI Sales Orchestrator (within C6) | 3-7+ years

LLM-driven sequencing, intent-based lead scoring, CRM-AI integrations. Bullet formula adds: Implemented AI tool X to automate Y resulting in Z.

Conversation AI Analyst (within C6) | 3-5+ years

Prompt-tuning, conversation-flow design, escalation rules, KPI-driven script-testing. Resume skews technical — tools section must include specific platforms.

AI Feature Launch Manager (within C7) | 5-8+ years

AI-explanation-level clarity, ethics narrative, regulatory readiness (EU AI Act, India DPDP). Summary must reference responsible AI or explainability.

Generative Content Orchestrator (within C7) | 3-7 years

Prompt engineering for content, brand-voice consistency across AI outputs, rights-management. Portfolio link mandatory.

AI Sourcing Analyst (within C8) | 3-7 years

AI sourcing tools, Talent Graph APIs, GDPR/DPDP data protection. Bullet formula adds: Used AI tool X to source Y candidates resulting in Z time-to-fill improvement.

FinOps Engineer (within C5) | 5-10+ years

Cloud-cost modelling, optimisation, policy-based guardrails. Bullet metric is always cost saved, not just infra built.

Prompt Engineer / AI Interaction Designer (within C1) | 3-7 years

LLM behaviour, user-intent mapping, conversational design. Resume bridges tech and design — include both code evidence and UX thinking evidence.

SECTION 3: LAYER 2 — ATS INTELLIGENCE

ATS systems do not reject candidates. They rank and deprioritise them. A resume that fails ATS scoring is not auto-rejected — it is ranked low enough that the recruiter never sees it. The problem is not rejection. It is invisibility.

> **Corrected Framing for ATS Optimisation**
> The widely cited '75% of resumes are auto-rejected by ATS' is substantially incorrect. Primary research (Enhancv 2025, n=25 recruiters) shows only 8% of ATS systems are configured for content-based auto-rejection. The real mechanism: recruiters review the top 15–25 results from 400–1,000 applications. ATS optimisation means ranking in the top 15–25 — not avoiding rejection.

3.1 ATS Detection Protocol

When a candidate provides a job link, the system detects the ATS from URL signature before any other processing occurs.

ATS Platform	URL Signature	Behaviour + Optimisation Priority
Greenhouse	`boards.greenhouse.io`	Keyword + structured. Prioritises exact match, section hierarchy. PDF parses well. Keyword density in first half of resume critical.
Lever	`jobs.lever.co`	Keyword + slight semantic. More forgiving of structure variation. Favours GitHub links and portfolio references. PDF preferred.
Workday	`myworkdayjobs.com`	Heavy filtering. Rigid section parsing. Must use standard headers: Education, Experience, Skills. Word .docx parses more reliably than PDF.
Taleo	`taleo.net`	Rigid keyword. Most sensitive to formatting. Absolute: no tables, no columns, no special characters in section headers. Plain text most reliable.
Ashby	`ashbyhq.com`	Slight semantic. Most modern. Recognises skill synonyms. Portfolio and GitHub carry weight. PDF preferred.
SmartRecruiters	`smartrecruiters.com`	Keyword + skills. Strong skills section weighting. Certifications parse well. Word preferred for enterprise applicants.
SAP SuccessFactors	`successfactors.com`	Enterprise-grade filtering. Exact tool names critical: 'Power BI' not 'PowerBI'. Section order matters. Word preferred.
Unknown ATS	`None detected`	Apply generic ATS-safe rules: single column, standard headers, no tables in body, PDF format. Assign Generic ATS Compatibility Score.

3.2 ATS Scoring Model

The Resume OS applies a standardised ATS scoring model to simulate how each variant will rank before output.

Keyword Match 35%	Exact and semantic match between JD keywords and resume content across Summary, Skills, and Experience sections.

Title Match 25%	Alignment between candidate's most recent title and the target role title. ATS systems weight current/last title heavily.
Experience Relevance 20%	Depth and recency of experience in the required functional area. Measured by density of relevant terminology in experience bullets.
Skills Section 10%	Presence of exact tool and skill names listed in the JD. Skills section must be structured as plain text, not icons or rated bars.
Parsing Gate Pass/Fail	Formatting check. If the resume fails parsing (tables in body, multi-column layout, embedded images, non-standard headers), all other scores are invalidated.

3.3 Universal ATS Formatting Rules (Hard Constraints)

These rules apply regardless of ATS type. Violation of any rule degrades the ATS score. Some violations fail the Parsing Gate entirely.

- Section headers must use standard names: 'Experience' not 'Where I've Worked', 'Skills' not 'My Toolkit', 'Education' not 'My Background'
- No tables in the resume body. Tables appear to be empty to most ATS parsers.
- No multi-column layouts for roles targeting Workday or Taleo. Greenhouse and Lever tolerate single visual columns inside a structured layout.
- No icons, rating bars, or graphical skill representations. These are invisible to parsers.
- Skills must appear as plain comma-separated text in a dedicated Skills section.
- File format: PDF for Lever, Greenhouse, Ashby. Word .docx for Workday, Taleo, SAP SuccessFactors.
- Exact tool names matter: 'Power BI' not 'PowerBI', 'Microsoft Excel' not 'Excel' for enterprise roles, 'Node.js' not 'NodeJS'
- GitHub profile, portfolio links, and LinkedIn URL must appear in the header contact block, not embedded in body text.

3.4 ATS Compatibility Score Calculation

The ATS Compatibility Score is calculated for each of the 7-10 output variants. It is reported as a percentage alongside each variant in the output package.

Score Interpretation

85-100%: Strong — likely to rank in top 15-25 results70-84%: Good — competitive, minor improvements possible55-69%: Fair — will appear but below top results; re-optimisation recommendedBelow 55%: Weak — will not surface in recruiter's shortlist; do not submit until re-optimised

SECTION 4: LAYER 3 — RECRUITER DECISION MODEL

ATS gets the resume seen. The recruiter decides in 6 seconds whether to read it. These are two separate problems. The recruiter is not reading — they are scanning. Resume OS optimises for both the scan and the read.

4.1 The 6-Second Scan Model

Eye-tracking research (The Ladders 2025) confirms that initial recruiter review of a resume averages 6-8 seconds. In that window, the recruiter checks four things in order:

1. CLARITY	Can I understand what this person does in 3 seconds? Cluttered formatting, unclear titles, or walls of text fail here immediately.
2. RELEVANCE	Is this person's background relevant to my open role? Title match, company recognition, and industry alignment are the signals.
3. PROOF	Is there evidence this person can do the job? Metrics, scale indicators, and named systems/tools are the proof nodes.
4. RISK	Is there anything that makes this a risky hire? Gaps, trajectory inconsistency, vague ownership, or unverifiable claims register here.

> **Recruiter Optimisation Rule**
> The resume must pass all four scan checks in the top half of page one. If clarity, relevance, proof, and risk signals are not readable in the top 40% of the document, the recruiter has already moved to the next candidate.

4.2 Risk Signal Register — Signals to Eliminate

Hiring is a risk decision. The recruiter is making a bet with their reputation. Every risk signal on the resume increases the probability of a pass. Resume OS scans for and eliminates all of the following before output.

Risk Signal	Why It Registers as Risk	Fix
Vague ownership ('Worked with the team to...')	Signals no personal accountability. Recruiter cannot determine what this person actually did.	Rewrite to: 'Owned X. Led Y. Designed Z.'
No metrics in experience section	Unquantified claims are unverifiable. Signals either low impact or low self-awareness.	Minimum 1 metric per every 2 bullets. Use scope as proxy if outcome metric unavailable.
Inconsistent trajectory (frequent short tenures without explanation)	Signals flight risk, performance issues, or inability to fit. Particularly penalised for L3+ roles.	Front-load longest tenures. Add brief context for transitions where relevant.
Over-generalisation (listed 25+ skills, claims 5 different specialisations)	Signals lack of real depth. Recruiters prefer a clear specialist over a vague generalist at mid-senior level.	Cluster skills into 2-3 domains. Choose the specialisation most relevant to the target role.
'Responsible for' language	Describes a job description, not an achievement. Signals task-taker mentality, not owner mentality.	Replace with action verbs: Led, Built, Designed, Reduced, Grew, Shipped, Defined.
Soft skill claims without evidence ('Excellent communicator')	Unverifiable. Recruiter has read 'excellent communicator' 400 times today. Signals low signal quality.	Remove entirely, or replace with evidence: 'Presented monthly reviews to C-suite of 8'

Risk Signal	Why It Registers as Risk	Fix
'Familiar with' or 'exposure to'	Signals inability to use in a job context. ATS may count it but recruiter filters it out.	Either demonstrate use with a project/outcome, or remove. No middle ground.
College projects at 5+ years of experience	Signals the professional experience is not strong enough to fill the resume. Creates recency confusion.	Remove. Replace with projects that demonstrate beyond job scope if needed.

SECTION 5: LAYER 4 — GAP ANALYSIS ENGINE

Gap analysis runs after Role Intelligence and ATS Intelligence are complete. It compares the structured signals extracted from the JD against the signals present in the candidate's master resume. The output is a gap scorecard that drives all optimisation decisions in Layer 5.

5.1 JD Signal Extraction

Before comparison, the JD is parsed for the following structured signal types:

- Required skills: Explicit skill mentions with 'required', 'must have', 'essential'
- Preferred skills: Explicit mentions with 'preferred', 'nice to have', 'bonus'
- Title signals: The exact job title and any mentioned comparable titles
- Level signals: Years of experience requested, seniority descriptors, reporting lines
- KPI signals: Metrics mentioned — quota, uptime, throughput, NPS, conversion rate, time-to-fill, etc.
- Industry signals: Domain mentioned — fintech, healthcare, SaaS, e-commerce, etc.
- Environment signals: Company stage, team size, growth rate, culture descriptors
- Action verbs: What the role requires the person to DO — verbs used in job responsibilities

5.2 Gap Scorecard — Output Format

The Gap Analysis produces a scorecard across five dimensions. This scorecard is used internally to prioritise which Optimization Engine sub-systems to apply most aggressively.

Dimension	What Is Measured	Threshold for Strong	Action If Below Threshold
Keyword Coverage	% of JD required keywords present in resume	70%+ required, 50%+ preferred	Keyword Engine: add missing keywords with evidence where possible; note gaps in strategy log
Title Alignment	Similarity between candidate's most recent title and target title	Exact match or 1-level adjacent	Title match: adjust section framing; note if title gap is structural — candidate must be informed
Role Match Score	Overlap of functional experience with JD requirements	75%+	Role-Shift Engine: reframe bullets to emphasise most relevant experience cluster
Leadership Signal	Presence of ownership, decision, and leadership language	Strong for L3+	Signal Amplification: replace participation language with ownership language throughout
Risk Level	Number of risk signals detected (from Section 4.2)	0 hard risks; max 2 soft	Risk elimination: remove or reframe all detected risk signals before output

Integrity Constraint

Gap Analysis may identify keywords or skills that are absent from the candidate's master resume. Resume OS may surface this information to the candidate as a self-improvement signal. It may NOT add missing skills or experiences to the resume. The output is always bound by what is in the master resume. Adding fabricated signals is a disqualification risk and a system integrity violation.

SECTION 6: LAYER 5 — OPTIMIZATION ENGINE

The Optimization Engine contains 8 sub-engines that run in sequence. Each sub-engine takes the gap scorecard and role DNA as inputs and applies targeted rewrites, restructuring, and signal amplification to the master resume content.

The master resume is never overwritten. All engines operate on a working copy. The original master is preserved.

SUB-ENGINE 5.1 — POSITIONING ENGINE

Purpose: Rewrite the professional summary or objective to create precise role-specific positioning. Every version must answer three questions explicitly:

Who you are	Your functional identity, specialisation, and years of relevant experience. Not generic — tied to the target role's function and level.
For whom	The type of organisation, team, or problem you are best suited to. Startup vs enterprise. Product vs service. B2B vs B2C. Match to environment axis.
Driving what outcome	The specific business or technical outcome you reliably produce. This is the closing proof statement. It must reference a real achievement or capability.

Summary length rules by level:

- Entry (L1-L2): 3 lines. Anchor: education + project/intern proof + role seeking.
- Mid (L3): 3-4 lines. Anchor: specialisation + ownership proof + directional ambition.
- Senior (L4): 4-5 lines. Anchor: domain authority + org impact + unique positioning.
- Lead/Principal (L5+): 5-6 lines. Anchor: systems built + org scale + point of view.

SUB-ENGINE 5.2 — KEYWORD ENGINE

Purpose: Maximise ATS keyword coverage without keyword stuffing. The Keyword Engine follows a distribution model — not concentration.

- Extract: Pull exact required and preferred keywords from the JD.
- Expand: Add documented semantic variants (e.g., 'machine learning' and 'ML' and 'AI models' — all three may be needed depending on ATS type).
- Distribute: Place keywords across three sections in order of ATS weighting:
  - Summary section: 2-4 core role keywords integrated naturally into positioning statements
  - Experience section: Exact keywords appear in bullet context with evidence attached
  - Skills section: Full list of relevant keywords in plain comma-separated format
- Verify: No keyword appears in Skills alone without evidence in Experience. Unsupported keyword claims are risk signals to human reviewers.

SUB-ENGINE 5.3 — BULLET ENGINE

Purpose: Rewrite every bullet in the experience section to follow the Role-calibrated bullet formula. This is the highest-impact sub-engine for recruiter optimisation.

Master bullet formula:

Formula

Action Verb + Scope + Impact + MetricExample: Designed (action) the event-driven notification system serving 800K users (scope) reducing support ticket volume (impact) by 34% in 3 months (metric)

Level-calibrated formula adjustments:

- L1-L2 (Entry): Built [what] using [tech/method] that [does what] for [who/scale]. Metric = system scale or quantified output.
- L3 (Mid): Owned/Led [system or process] that [outcome] impacting [scope + metric]. Metric = business result.
- L4 (Senior): [Action at org level] — [decision made] — [business/scale outcome + metric]. Metric = org-level change.
- L5+ (Lead): [System/team built] + [org transformation] + [business impact metric]. Individual contribution language removed.

Cluster-specific bullet formula variants:

C1 — Engineering
Formula: Action + Stack + Scale + Metric
Example: Designed the auth service (Node.js + Redis) handling 2M daily logins; reduced latency by 62%.

C2 — Data/BA
Formula: Analysis Type + Tool + Insight + Decision + Result
Example: Built CLV model (Python/sklearn) segmenting 180K users; informed ₹3.2Cr quarterly budget reallocation.

C3 — DS/MLE
Formula: Model + Dataset + Framework + Metric vs Baseline
Example: Trained XGBoost classifier on 85K-row dataset; AUC-ROC 0.89 vs 0.74 logistic baseline.

C4 — Product
Formula: Product Area + Decision Made + User/Business Outcome + Metric
Example: Owned integrations marketplace; contributed to 34% increase in enterprise contract renewals.

C5 — DevOps
Formula: Infrastructure + Tool + Automation/Reliability Outcome + Metric
Example: Migrated 18 services to GKE; reduced infrastructure cost ₹22L/year through autoscaling.

C6 — Sales
Formula: Activity/Deal + Scale + Pipeline/Revenue Outcome + Quota %
Example: Closed ₹2.8Cr new ARR in FY24 — 131% of quota; largest deal ₹48L multi-year contract.

C7 — Marketing
Formula: Channel + Budget + Metric vs Benchmark
Example: Managed ₹1.2Cr/month paid media; delivered 2,400 MQLs at ₹3,330 CAC vs ₹4,800 target.

C8 — HR/TA
Formula: Process + Volume + Outcome Metric
Example: Coordinated recruitment for 12 roles; average time-to-fill 28 days vs 45-day team benchmark.

C9 — Ops/Consulting
Formula: Situation + Action + Quantified Result (SAR)
Example: Redesigned warehouse picking; reduced errors 41%, increased throughput 800 to 1,150 units/day.

C10 — Design
Formula: Problem + Research Method + Design Decision + User/Business Outcome
Example: Redesigned checkout: 8 interviews, 3 iterations, usability test n=10 → abandonment 67% to 19%.

SUB-ENGINE 5.4 — ROLE-SHIFT ENGINE

Purpose: Ensure every bullet, section, and framing decision matches the target level's hiring expectations. This prevents the most common failure mode: a 5-year candidate applying to L4 roles with L3-language resumes.

The engine scans all content and flags every instance of language that signals a level lower than the target. It then rewrites to level-appropriate language without changing the underlying facts.

L1 to L3 shift	'I helped build a dashboard' → 'Owned the analytics dashboard design and weekly distribution to 12 stakeholders'
L2 to L4 shift	'Worked on microservices migration' → 'Led the microservices decomposition initiative across 3 backend services; designed the event-driven architecture and managed 2-engineer pod'
L3 to L5 shift	'Led a team of 4 engineers' → 'Scaled and managed a 4-person backend engineering team; defined hiring criteria, conducted all technical interviews, and established team-level engineering standards'

SUB-ENGINE 5.5 — ENVIRONMENT ADAPTATION ENGINE

Purpose: Adjust tone and framing to match the cultural hiring bias of the target organisation type. Same facts, different emphasis.

Startup Mode	Amplify: 'built from scratch', 'owned end-to-end', 'zero to one', 'shipped in X weeks', 'operated without process'. Reduce: governance language, approval chains, committee decisions.
SME Mode	Amplify: 'cross-functional', 'wore multiple hats', 'supported both X and Y', 'pragmatic solutions'. Balance specialist depth with versatility signals.
MNC Mode	Amplify: 'managed across X regions', 'compliant with Y framework', 'aligned with Z stakeholders', 'scaled within structured program'. Reduce: chaos, solo execution, no-process language.

SUB-ENGINE 5.6 — SIGNAL AMPLIFICATION ENGINE

Purpose: Surface the three highest-value signal types — metrics, leadership verbs, scale indicators — and increase their density in the top half of the resume.

- Metrics visibility: Every experience section must have at least one metric in the first two bullets. If the first visible proof is bullet 4, amplify bullets 1-2 first.
- Leadership verbs: Replace all passive and participation verbs with the strongest accurate ownership verb. Hierarchy: Architected > Designed > Led > Built > Delivered > Managed > Contributed.
- Scale indicators: Add scale where available — team size, user count, transaction volume, budget managed, geographic scope. Scale without metric is better than neither.

SUB-ENGINE 5.7 — RECENCY WEIGHTING ENGINE

Purpose: Weight the density and depth of content toward the most recent and most relevant experience. Older experience compresses progressively.

Current / most recent role	4-6 bullets. Full bullet formula applied. Maximum detail.

1 role back	3-4 bullets. Full formula applied. Slightly compressed.
2 roles back	2-3 bullets. Key achievements only. No task descriptions.
3+ roles back	1-2 bullets or company + title + dates only. No detail unless directly relevant to target role.
Education	Moves progressively toward bottom. At 7+ years: 2 lines only — degree + institution + year. No coursework, no GPA unless specifically relevant.
Exception	If a past role is MORE relevant to the target JD than recent roles, the Recency Weighting Engine notes this in the strategy log and recommends a variant that surfaces that role more prominently.

SUB-ENGINE 5.8 — STRUCTURE ENGINE

Purpose: Determine the optimal section order for each variant based on role cluster, experience level, and ATS target. Structure is not cosmetic — it determines what the recruiter sees in the first 6 seconds.

Structure rules by experience level:

Entry Level (L1-L2)
Order: Education → Skills → Projects → Experience → Certifications
Education is still the primary credibility signal. Projects and personal builds substitute for work experience.

Mid Level (L3)
Order: Summary → Skills → Experience → Projects (optional) → Education
Work experience is now primary. Summary does positioning work. Education recedes.

Senior Level (L4)
Order: Executive Summary → Core Competencies → Experience → Technical Proficiencies → Education
Individual skills list replaced by competency domains. Tools appear only at bottom for ATS.

Lead / Principal (L5+)
Order: Leadership Thesis → Scope Markers → Experience → Notable Contributions → Speaking/Publications → Education
Resume speaks to the business. Scope markers (team, budget, org) are immediately visible.

Cluster-specific structure overrides (applied on top of level rules):

- C3 DS/MLE Entry: Education surfaces first — institution and CGPA are primary signals in research-adjacent roles. Publications appear above projects.
- C5 DevOps/Cloud Entry: Skills block moves above Education — tool certifications (AWS, GCP, Azure) are the primary differentiator and must be front-loaded.
- C10 Design all levels: Portfolio link appears as the first element in the contact header. The resume is a door. The portfolio is the room.
- C6 Sales mid-senior: Key Metrics block (quota %, ACV, ARR managed) appears as a dedicated section immediately after the summary. Numbers must be visible without reading any bullets.
- C3 Senior DS: Structure diverges from MLE — Research Thesis replaces Executive Summary. Model Portfolio replaces Core Competencies. Publications appear above Tools.

SECTION 7: LAYER 6 — OUTPUT ENGINE

The Output Engine generates 7-10 resume variants from the optimised master content, plus a scoring package and strategy log. All variants contain identical factual content from the master resume. They differ in structure, section order, tone, and the specific signals each variant foregrounds.

7.1 The 7-10 Variant Model — Why Variants Exist

A single recruiter reviewing the same role will spend 6-8 seconds on a resume. Different recruiters at different companies have different trained scan patterns. Some read top-to-bottom. Some scan for titles first. Some look for metrics immediately. Some respond to narrative summaries. Some respond to skill-dense layouts.

Variants are not cosmetic alternatives. They are strategic signal rearrangements. The candidate chooses the variant whose structure best matches the recruiter context — or uses the default highest-scoring variant for blind applications.

7.2 The 10 Base Variants — Structure and Target Context

All 10 variants apply to every role cluster. The Role DNA and ATS Detection layers adjust the content of each variant — the variant templates define the structure logic.

V1 — The Signal Stack

Target: Tech roles: Software Engineer, DevOps, ML Engineer at product companies

ATS: Greenhouse, Lever

Structure: Skills → Projects → Experience → Education

Logic: Front-loads the keyword block for ATS. Leads with tech depth. Favoured by engineering managers who scan for stack first.

Tone: Technical, direct, evidence-led. Minimal narrative. Every bullet is a proof point.

V2 — The Outcome Ledger

Target: PM, Sales, Marketing, Business roles where outcomes are the primary hire signal

ATS: Greenhouse, Workday

Structure: Summary → Metrics Block → Experience → Skills → Education

Logic: Surfaces numbers in the first quarter of the resume. Designed for roles where revenue, growth, or product outcomes determine the hire decision.

Tone: Results-first. Assertive. Business language throughout.

V3 — The Authority Frame

Target: Enterprise/MNC roles: DA, BA, Ops, Consulting, HR in structured organisations

ATS: Workday, SAP SuccessFactors

Structure: Summary → Tools & Certifications → Experience → Education

Logic: Leads with structured credentials and exact tool names. Optimised for Workday/SAP heavy-filtering. Formal, process-fluent framing.

Tone: Formal, process-led, governance-aware. MNC language throughout.

V4 — The Leadership Thesis

Target: Senior and Lead roles: L4-L5 in any function. Manager, Director, VP track.

ATS: Any ATS

Structure: Executive Summary → Team/Scale Markers → Experience → Core Competencies → Education

Logic: Org-level positioning from line one. Scale markers (team, budget, org size) visible without reading bullets. Designed for executives reviewing senior hires.

Tone: Executive. Strategic. Org-level impact language. Individual contribution language minimised.

V5 — The Proof Sheet

Target: Research, ML, Data Science, Academic/industry hybrid roles

ATS: Lever, Ashby

Structure: Research Interests → Publications/Papers → Projects → Skills → Education (top)

Logic: Inverts the standard structure — Education and research output come first. Designed for roles where academic credibility anchors the entire profile.

Tone: Academic precision. Methodology-led. Metric comparisons vs baseline are mandatory.

V6 — The Problem-Solver

Target: Ops, Consulting, Supply Chain, Strategy, and cross-functional execution roles

ATS: Workday, SmartRecruiters

Structure: Summary → Core Competencies Matrix → Experience (SAR format) → Certifications → Education

Logic: SAR (Situation-Action-Result) bullet structure mirrors case interview format — the primary evaluation lens for consulting and ops roles. Competency matrix replaces skills list.

Tone: Structured, analytical, client-facing. Problem-solution language throughout.

V7 — The Portfolio Lead

Target: UX/UI Design, Product Design, Creative and visual roles

ATS: Lever, Ashby, Greenhouse

Structure: Portfolio Link (hero) → Summary → Case Studies → Tools → Education

Logic: The resume is a door. Portfolio link is hero placement. Case studies replace experience bullets — each follows Problem → Process → Decision → Outcome. Minimises text, maximises proof by link.

Tone: Process-led, user-centric. Design thinking language throughout. Minimal prose.

V8 — The Versatility Map

Target: SME and multi-role positions requiring breadth over depth. Generalist track.

ATS: Greenhouse, Workday

Structure: Summary → Cross-functional Skills → Experience (breadth-emphasised) → Education

Logic: Surfaces range explicitly. Reframes each role to show breadth of contribution — 'supported both X and Y', 'owned across ops and marketing'. For roles where versatility is the primary hiring signal.

Tone: Pragmatic, adaptable. Cross-functional language. 'Wore multiple hats' signals are amplified.

V9 — The Domain Expert

Target: Deep specialist roles: Senior+ in a single domain. Fintech, healthcare, SaaS vertical.

ATS: Greenhouse, Lever, Workday

Structure: Domain Summary → Industry Credentials → Experience (depth-led) → Publications/Certifications → Education

Logic: Positions the candidate as THE expert in one domain. Suitable for roles where domain knowledge is the hire decision — e.g., BFSI risk, healthcare compliance, SaaS product. All content filtered for domain relevance.

Tone: Authority. Deep expertise language. Domain-specific vocabulary throughout.

V10 — The Transition Narrative

Target: Career changers, re-entrants, pivots between functions or industries

ATS: Greenhouse, Lever

Structure: Transition Summary → Transferable Skills → Relevant Experience → New-Domain Projects → Education

Logic: Leads with a positioning summary that explicitly names the transition and makes the case for transferability. Reframes past experience through the lens of the target role. New-domain projects (personal, freelance, or formal) surface prominently.

Tone: Forward-looking. Explicit about the pivot. Evidence of active preparation for new domain.

7.3 Scoring Package — Per Variant

Each of the 7-10 output variants receives a score on four dimensions. Scores are reported alongside each variant to help the candidate select the best option for their target application.

ATS Compatibility Score (0-100%)	Simulated ATS rank based on keyword match (35%), title alignment (25%), experience relevance (20%), skills (10%), parsing (gate). ATS-specific model applied where ATS is detected.
Recruiter Scan Score (0-100%)	Clarity, relevance, proof, and risk signals evaluated for the top 40% of the resume. Measures how well the 6-second scan check passes.
Signal Density Score (0-100%)	Metrics per bullet, leadership verb ratio, scale indicator frequency. Measures how much verifiable proof is packed into the resume.
Role Match Score (0-100%)	Overlap between candidate experience signals and the specific role requirements from the JD. Combines keyword coverage and functional relevance.

7.4 Strategy Log — Required Output

Every Resume OS run produces a strategy log alongside the resume variants. The strategy log is a structured record of every optimisation decision made. It serves three purposes:

- Transparency: The candidate understands exactly what changed and why. Nothing is changed without documentation.
- Learning input: Strategy logs feed the continuous learning system (see Section 8). Patterns across logs build the compound intelligence layer.
- Candidate coaching: The log often reveals what to develop — if keyword gaps are structural (missing skills, not missing words), the candidate is informed.

Strategy Log Format:

Strategy Log Structure

INPUT SUMMARYRole DNA detected: [Function + Level + Environment + KPIs]ATS detected: [ATS name or 'Unknown — generic rules applied']Gap Scorecard: [Keyword Coverage % | Title Alignment | Role Match | Leadership Signal | Risk Level]CHANGES MADE[Change 1]: [What changed] — [Why] — [Expected impact][Change 2]: ...KEYWORD GAPS NOT ADDRESSABLE[Keyword]: Not present in master resume. Candidate may want to develop: [learning path]VARIANT RECOMMENDATIONHighest ATS score: V[X] — [score]%Highest recruiter score: V[X] — [score]%Recommended for this role: V[X] — [rationale]

SECTION 8: CLUSTER-LEVEL TEMPLATE SPECIFICATIONS

This section contains the complete template specification for each of the 10 role clusters, across all experience bands. Each specification defines the exact structure order, positioning engine output, bullet formula, ATS rules, and risk signals for that cluster-level combination.

These specifications are the source of truth for all AI-generated resume variants. When generating output for a candidate, the system selects the matching Cluster + Experience Band specification and applies it.

CLUSTER 1: SOFTWARE ENGINEER / FULL-STACK / BACKEND / FRONTEND

ATS Priority	Greenhouse (keyword + structured), Lever, Workday (enterprise)
Core Keywords	system design, REST API, microservices, cloud deployment, CI/CD, performance optimisation, scalability, distributed systems
Bullet Formula	Action + Stack + Scale/Outcome + Metric

C1-E | Entry Level | 0-2 years

Structure: Education → Skills → Projects → Internship/Experience → Certifications

Rationale: Projects ARE experience at this level. ATS keyword block immediately after education front-loads match. Education stays top because GPA and institution are legitimate signals.

Positioning: Computer Science / Engineering graduate with hands-on project experience in [primary stack]. Designed and deployed [X] production-grade systems independently. Seeking to contribute to [startup product / enterprise engineering team] from day one.

Risk flags: 'Familiar with' language; generic project descriptions; 20+ skills without project evidence; missing GitHub link.

C1-M | Mid Level | 3-6 years

Structure: Summary → Skills (condensed, 8-12 items) → Experience → Projects (only if beyond job scope) → Education

Rationale: Work experience is now primary. Summary does positioning. Skills stay for ATS but become selective. College projects disappear unless they are genuinely beyond job scope.

Positioning: [X] years building [specific system type] in [primary stack]. Own the [specific service/system] at [company type] — [scale indicator]. Looking to take on [system design / architectural / team] responsibility in [environment target].

Risk flags: College projects still prominent; participation language; inconsistent title trajectory; skills list unchanged from entry-level.

C1-S | Senior Level | 7-12 years

Structure: Executive Summary → Core Competencies (capability domains, not tools) → Experience → Technical Proficiencies → Education (2 lines)

Rationale: Recruiter at senior level asks: 'Can this person solve our hard technical problem?' Summary must answer. Core Competencies replaces skills list. Tools appear at bottom for ATS only.

Positioning: Senior [function] Engineer with [X] years building [domain] systems across [industries]. [Specific org-level achievement — architecture decision, scale reached, team led]. Deep expertise in [stack]. Known for [unique positioning statement].

Risk flags: Still writing L3-style task bullets; no architecture or org-level decisions visible; skill list unchanged from mid-level; no scale indicators.

C1-L | Lead / Principal / Staff | 12+ years

Structure: Leadership Thesis → Scope Markers → Experience → Notable Technical Contributions → Speaking/Publications → Education

Rationale: Resume talks to the business. Code mention is minimal. The question being answered is: 'What engineering systems and culture did this person create?'

Positioning: Principal/Staff Engineer with [X] years across [domains]. Built and led [team/platform] from [X] to [Y]. [Technical author / Conference speaker / Open source maintainer]. Currently seeking [scope] role where technical strategy shapes product direction.

CLUSTER 2: DATA ANALYST + BUSINESS ANALYST

ATS Priority	Workday (dominant), SmartRecruiters, SAP SuccessFactors
Core Keywords	SQL, data analysis, dashboards, Power BI, Tableau, Excel, stakeholder reporting, requirements gathering, business insights, cohort analysis
Bullet Formula	Analysis Type + Tool + Insight Found + Business Decision/Action + Measurable Result
ATS Critical Rule	SQL must appear as standalone keyword. 'Power BI' not 'PowerBI'. 'Microsoft Excel' not 'Excel' for enterprise roles.

C2-E | Entry Level | 0-2 years

Structure: Education → Technical Skills → Projects/Internship → Certifications → Extracurriculars

Rationale: Quantitative degree (Math/Stats/CS/Commerce) is the primary credibility signal at entry. Certifications (Google Data Analytics, SQL) unusually important — place above projects.

Positioning: [Degree] graduate with hands-on experience in SQL, Python (pandas), and [BI tool]. Completed [internship/project] building [specific output] used by [stakeholder]. Comfortable translating raw data into structured business recommendations.

C2-M | Mid Level | 3-6 years

Structure: Summary → Skills/Tools → Experience → Key Projects (if distinct from job) → Education

Positioning: Data/Business Analyst with [X] years turning complex datasets into [business decisions / product insights] across [industries]. Own end-to-end analytics for [team] — from SQL extraction to [tool] dashboard to stakeholder recommendation. Strong in [specialisation].

C2-S | Senior Level | 7-12 years

Structure: Summary → Core Competencies (analytical domains) → Experience → Tools (bottom) → Education

Shift: From analysis execution to analytics strategy. Define what to measure. Build analytical culture. Influence roadmap with data.

Positioning: Senior Data/Business Analyst with [X] years across [domains]. Built the analytics function at [company type] — defined KPI frameworks, implemented tracking infrastructure, [team growth]. Translated [X]+ analytical recommendations into product and marketing decisions.

CLUSTER 3: DATA SCIENTIST + ML ENGINEER

ATS Priority	Lever (semantic), Ashby, Greenhouse

Core Keywords	machine learning, Python, model training, feature engineering, model deployment, experimentation, deep learning, TensorFlow, PyTorch, MLOps
Bullet Formula (Mandatory)	Model Type + Dataset/Domain + Framework + Metric vs Baseline. Metric without baseline is not acceptable.
Senior Split	DS and MLE diverge structurally at 7+ years. Treat as separate templates from Senior band.

C3-E | Entry Level | 0-2 years

Structure: Education (CGPA matters — include it) → Research/Publications (under review counts) → Projects → Skills → Certifications

Rationale: This is the only cluster where Education surfaces first regardless of level. Institution and CGPA are primary signals for research-adjacent roles.

Critical rule: Every ML project bullet must state: model type, dataset size, framework, metric achieved, metric vs baseline. No exceptions.

C3-M | Mid Level | 3-6 years

Split treatment begins here. DS and MLE have distinct summaries and bullet framing:

DS Summary: [X] years specialising in [NLP / CV / RecSys / Pricing]. Shipped [X] models to production collectively driving [business outcome]. Strong in experiment design: [X] concurrent A/B tests per quarter with rigorous statistical controls.

MLE Summary: [X] years building and deploying ML systems at scale. Own the model serving infrastructure processing [X] predictions/day on [platform]. Reduced [cycle time / latency / cost] by [metric] through [method].

C3-S | Senior Level | 7-12 years

DS Structure: Research Thesis → Model Portfolio (3-5 models + business outcomes) → Experience → Publications → Tools

MLE Structure: Platform Engineering Thesis → Systems Owned (scale, reliability, impact) → Experience → Tools → Education

CLUSTER 4: PRODUCT MANAGER + PRODUCT ANALYST

ATS Priority	Greenhouse, Lever, Workday
Core Keywords	product roadmap, user research, A/B testing, KPIs, cross-functional, stakeholder alignment, go-to-market, product strategy, PRD
Bullet Formula	Product Area + Decision Made + User/Business Outcome + Metric
Entry Critical Rule	Resume must show evidence of product THINKING, not just execution. Signal: Did this person make a product decision and measure its outcome?

C4-E | Entry Level | 0-2 years | APM / Junior PM / Product Analyst

Structure: Summary → Skills (tools: Jira, Figma, SQL, Amplitude, Mixpanel) → Projects/Internship (product decisions, not features) → Education → Extracurriculars

C4-M | Mid Level | 3-6 years

Structure: Summary → Core Skills → Experience (outcome-led) → Education

Positioning: Product Manager with [X] years building [B2B/B2C, SaaS/marketplace] products in [domain]. Own [product surface] — [X] major releases shipped in [timeframe], [metric outcome]. Strong in [0-to-1 / data-informed / platform] product development.

C4-S | Senior Level | 7-12 years

Structure: Executive Summary → Product Portfolio → Experience (strategy + org-level decisions) → Education

Shift: From feature ownership to platform/portfolio ownership and business strategy. Every bullet must reference a strategic decision or business outcome — not a feature shipped.

CLUSTER 5: DEVOPS + CLOUD ENGINEER + SRE

ATS Priority	Greenhouse, Lever, Workday
Core Keywords	Kubernetes, Docker, Terraform, CI/CD, AWS/GCP/Azure, infrastructure as code, monitoring, reliability, SLO, SLI, automation, GitOps
Bullet Formula	Infrastructure + Tool + Automation/Reliability Outcome + Cost/Reliability Metric
Certifications Rule	Cloud certifications (AWS, GCP, Azure) are genuine differentiators at entry level. Must appear prominently — not buried. At mid-senior, certifications move below experience but remain.

C5-E | Entry Level | 0-2 years

Structure: Skills (cloud platforms, tools) → Education → Projects (infrastructure built, not just used) → Certifications → Internship

Rationale: Skills block moves above Education in this cluster. Cloud certs are the primary differentiator at entry. Projects must demonstrate infrastructure built — not software deployed on existing infra.

C5-M | Mid Level | 3-6 years

Structure: Summary → Core Stack (specific: AWS EKS, Terraform, ArgoCD, Datadog) → Experience → Certifications → Education

Positioning: DevOps/Cloud/SRE Engineer with [X] years automating infrastructure and deployment workflows across [industries]. Own the [Kubernetes cluster / cloud infra / reliability function] running [X] services on [platform]. [Key metric: deployment frequency, cost reduction, uptime].

C5-S | Senior Level | 7-12 years | SRE Split

Senior DevOps/Cloud: Platform strategy thesis. Infrastructure portfolio owned (scale, services, regions). Org-scale automation and cost outcomes.

SRE adjustment: Emphasise error budgets, SLO/SLI definition, incident management programs, chaos engineering. Bullet metric shifts from 'built X' to 'achieved X uptime / reduced MTTR by Y'.

CLUSTER 6: SALES + BUSINESS DEVELOPMENT

ATS Priority	Salesforce-native ATS, Greenhouse, Workday
Core Keywords	pipeline, quota attainment, ARR, MRR, prospecting, outbound, account management, deal closure, revenue, Salesforce, HubSpot, CAC, ACV
Bullet Formula	Activity/Deal + Scale + Pipeline/Revenue Outcome + Quota %

HARD RULE	Every bullet must contain a number. Sales resumes without metrics are eliminated in the first pass. No exceptions.

C6-E | Entry Level | 0-2 years | SDR, BDR, Inside Sales

Structure: Summary (hunter identity) → Experience (all metrics) → Skills (CRM tools) → Education

Positioning: SDR/BDR with [X] months of outbound sales experience in [B2B SaaS / industry]. Consistently [above/at] quota — [X%] average over last [X] quarters. Strong in [cold email / LinkedIn / cold calling]. Looking to move into [closing role / senior SDR / AE].

C6-M | Mid Level | 3-6 years | AE, Sr. SDR, Account Manager

Structure: Summary → Key Metrics Block (quota %, ACV, total ARR, industries — visible immediately) → Experience → Tools → Education

Key Metrics Block is mandatory at this level. Format: Quota Attainment: [X%] avg (last [X] years) | ACV: [range] | Total ARR Managed: [₹/$/total] | Industries: [list]

C6-S | Senior / Enterprise | 7-12 years

Structure: Executive Summary → Career Metrics (total TCV closed, largest deal, team/territory) → Experience → Education

Shift: Strategic account management, team leadership, market development. Every bullet references deal complexity, stakeholder level (C-suite), or team outcome.

Emerging variants: AI Sales Orchestrator and Conversation AI Analyst — add AI tool stack to Skills section and include one bullet demonstrating AI-augmented workflow improvement with metric.

CLUSTER 7: MARKETING + GROWTH + CONTENT

ATS Priority	Greenhouse, Workday, Lever
Core Keywords	campaign management, performance marketing, SEO, SEM, email marketing, content strategy, growth experiments, CAC, ROAS, conversion rate, funnel, MQL, attribution
Critical Split	Performance/Growth (metrics-heavy, analytical) vs Brand/Content (narrative-heavy, portfolio-led). Tone and structure differ by track within levels.
Performance Bullet Formula	Channel + Budget + Metric vs Benchmark
Content Bullet Formula	Content Type + Audience + Reach/Engagement Metric + Business Outcome

C7-E | Entry Level | 0-2 years

Structure: Summary → Skills (tools: Google Ads, Meta Ads Manager, HubSpot, SEMrush) → Experience/Internship → Education → Portfolio link (content/brand track only)

C7-M | Mid Level | 3-6 years

Structure: Summary → Core Metrics (CAC, ROAS, MQLs — visible block) → Experience → Tools → Education

Positioning (Performance): Performance Marketing Manager with [X] years owning paid acquisition across [channels]. Manage [₹/$/budget]. [Key metric achievement — CAC reduction, ROAS improvement, MQL growth].

Positioning (Growth): Growth Marketer with [X] years running experiment-driven [acquisition/retention] programs. Owned full funnel. [Key metric outcome].

Emerging variants: AI Feature Launch Manager and Generative Content Orchestrator — summary must reference AI tool fluency and one bullet must demonstrate AI-augmented campaign or content outcome.

C7-S | Senior Level | 7-12 years

Structure: Summary (marketing leadership + revenue contribution) → Core Metrics Block (budget, pipeline %, team) → Experience → Education

Shift: From channel execution to demand generation strategy, team building, and revenue contribution. Budget managed and pipeline contribution percentage are mandatory scope markers.

CLUSTER 8: HR + TALENT ACQUISITION + PEOPLE OPS

ATS Priority	Workday (dominant), SAP SuccessFactors, Greenhouse
Core Keywords	hiring, recruitment, talent pipeline, onboarding, HRIS, employee engagement, performance management, sourcing, offer management, headcount planning, time-to-fill, cost-per-hire
Bullet Formula	Process + Volume + Outcome Metric
TA vs HRBP Split	TA track: hiring metrics (hires made, TTF, CPH, offer acceptance). HRBP track: org health metrics (retention, performance cycle completion, reorg outcomes).

C8-E | Entry Level | 0-2 years | HR Executive, Recruiter, TA Coordinator

Structure: Summary → Skills (HRIS: Workday, Darwinbox; ATS: Greenhouse; sourcing: LinkedIn Recruiter) → Experience → Education

C8-M | Mid Level | 3-6 years | Senior Recruiter, HRBP, TA Lead

Structure: Summary → Core Metrics Block (hires made, TTF, CPH, retention rates) → Experience → Tools/HRIS → Education

TA Positioning: Senior Technical/Functional Recruiter with [X] years hiring [profiles] at [company type]. Hired [X]+ professionals across [X] companies. Avg TTF: [X] days for IC, [X] days for senior+. [X]% hires from outbound sourcing.

HRBP Positioning: HR Business Partner with [X] years partnering with [function] organisations of [X]-[Y] employees. Supported [reorgs, PIPs, programs]. Known for translating business problems into people solutions.

C8-S | Senior / TA Leader / CHRO track | 7-12 years

Structure: Summary (people strategy + org contribution) → Scope Markers (headcount, team, budget) → Experience → Education

Emerging variants: AI Sourcing Analyst — add AI sourcing tools to skills and include bullet on AI-augmented pipeline outcome (time, volume, quality improvement).

CLUSTER 9: OPERATIONS + SUPPLY CHAIN + CONSULTING

ATS Priority	Workday, SAP SuccessFactors, SmartRecruiters

Core Keywords	process improvement, cost reduction, stakeholder management, project management, SOP, cross-functional, efficiency, vendor management, Six Sigma, lean, P&L
Bullet Formula (Mandatory)	SAR — Situation (what was broken) + Action (what I did) + Result (quantified improvement). Every bullet reads like a mini-case.
Certifications	PMP, Lean Six Sigma Green/Black Belt carry weight at all levels in this cluster. Place above Education at entry-mid, below Experience at senior.

C9-E | Entry Level | 0-2 years | Ops Analyst, Junior Consultant, SCM Analyst

Structure: Summary → Skills (Excel, SQL, PowerPoint, Visio, ERP basics) → Experience/Internship → Education → Certifications

C9-M | Mid Level | 3-6 years

Structure: Summary → Core Competencies (process improvement, vendor management, project delivery) → Experience (SAR) → Certifications → Education

Positioning: Operations/Consulting [title] with [X] years improving [process/supply chain/cross-functional] performance in [industry]. Track record of measurable cost reduction: [₹/$ savings] over [period]. Strong in [process redesign / vendor negotiation / project management].

C9-S | Senior Level | 7-12 years

Structure: Summary → Scope Markers (P&L, team, geography, budget) → Experience (org transformation) → Education/Certifications

Shift: From process execution to transformation strategy. Every bullet references org-level change, program scope, or cumulative financial impact.

CLUSTER 10: UX/UI DESIGNER + PRODUCT DESIGNER

ATS Priority	Lever, Greenhouse, Ashby (semantic-aware systems dominate for creative roles)
Core Keywords	UX design, user research, prototyping, Figma, wireframing, design system, usability testing, information architecture, interaction design
ABSOLUTE RULE	Portfolio link is mandatory at the top of every resume variant. The resume is a door. The portfolio is the room.
Bullet Formula	Problem Discovered + Research Method + Design Decision + User/Business Outcome

C10-E | Entry Level | 0-2 years

Structure: Portfolio link (hero) → Summary → Skills → Projects/Internship (process-led, not output-led) → Education

Rationale: Every bullet follows: Discovered problem → Research method → Designed solution → Tested with X users → Outcome metric.

C10-M | Mid Level | 3-6 years

Structure: Portfolio link → Summary → Case Studies (3 max) → Skills/Tools → Education

Critical shift: Move from 'projects' to 'case studies'. Each case study: Problem (1 sentence) → Process (research → design → test) → Decision (key design decision + rationale) → Outcome (task completion / engagement / NPS / support ticket reduction).

C10-S | Senior / Lead Designer | 7-12 years

Structure: Portfolio link → Design Leadership Thesis → Design Systems/Platform Work → Experience → Tools → Education

Positioning: Senior Product/UX Designer and design leader with [X] years building user experiences for [domains]. Founded/scaled design system to [X]+ components used across [X] product squads. Managed team of [X] designers. Design philosophy: [one sentence POV].

SECTION 9: CONTINUOUS LEARNING AND COMPOUNDING INTELLIGENCE

Resume OS is not a static rule system. It is designed to compound intelligence with every application cycle. This section defines the learning architecture — what is logged, what is analysed, how the system improves, and what signals feed back into the system.

> **The Compound Intelligence Principle**
> Every resume generated is a data point. Every application submitted is an experiment. Every outcome — interview received, no response, offer extended — is signal. Resume OS logs all of these and uses pattern analysis to improve future optimisations. Over time, the system builds a dataset of what works for which role type, at which level, at which company type, through which ATS.

9.1 What Is Logged Per Application

- Role DNA classification: Function + Level + Environment + KPIs
- ATS detected and formatting rules applied
- Gap scorecard: keyword coverage %, title alignment, role match, leadership signal, risk level
- Variant selected by candidate: which of the 7-10 variants was submitted
- Changes made: full strategy log with before/after content for every modified element
- Outcome (candidate-reported): No response / Recruiter screen / Technical screen / Final round / Offer / Decline
- Time to response (where tracked): days from submission to first recruiter contact

9.2 Pattern Analysis Layer — What the System Learns

After a minimum of 50 applications per cluster-level combination, the system runs pattern analysis to detect:

Variant performance by cluster	Which structural variant (V1-V10) produces the highest response rate for each role cluster. This becomes the default recommended variant for new applications in that cluster.
Keyword effectiveness	Which JD keywords, when present in the resume, correlate with higher response rates. Used to prioritise keyword placement in future optimisations.
Summary length optimality	Whether 3-line or 4-line summaries produce higher responses at each experience band. Applied as an updated rule in the Structure Engine.
Bullet density sweet spot	The number of bullets per role that correlates with highest response. Used to calibrate Recency Weighting Engine outputs.
ATS-specific optimisation drift	Whether specific ATS systems have changed their parsing behaviour — detected by outcome patterns diverging from predicted ATS scores.
Environment signal effectiveness	Whether startup-framed vs MNC-framed language produces different outcomes for the same role type. Refines Environment Adaptation Engine.

9.3 Learning Inputs — What Improves the System

The following inputs feed the learning layer. Each input type and its contribution is defined:

Input Type	Source	How It Improves the System
Application outcomes	Candidate-reported result after each application	Refines variant recommendation model; improves keyword priority weighting
Recruiter feedback	Feedback received at recruiter screen stage (when shared)	Identifies signal gaps not caught by gap analysis; adds new risk signals to register
JD corpus updates	New JDs processed across all clusters and levels	Updates keyword lists; detects emerging role variants; adds new cluster splits when divergence threshold reached
ATS behaviour changes	Outcome divergence from predicted ATS score	Triggers ATS Knowledge Base update; adjusts formatting rules for affected ATS platform
New role variant detection	When a JD cannot be classified into existing clusters with >70% confidence	Flags for human review; may generate new cluster or sub-cluster if pattern repeats
Candidate skill profile updates	When candidate adds new experience, certification, or project to master resume	Re-runs gap analysis across all active JDs; may upgrade variant recommendation

9.4 Integrity Guardrails — What the System Cannot Learn to Do

The learning layer is bounded by hard integrity constraints that cannot be overridden by any pattern, however strong.

Cannot learn to fabricate
No outcome data, however compelling, may cause the system to add experience, skills, metrics, or claims not present in the master resume.

Cannot learn to misrepresent level
No pattern may cause the system to claim a seniority level inconsistent with the actual years and evidence in the master resume.

Cannot learn to bypass ATS rules on safety
Formatting rules that protect parsing integrity cannot be overridden even if a specific variant with table elements shows a higher response rate in one data subset.

Cannot override candidate review
All AI-generated suggestions remain recommendations. The candidate reviews and approves all changes. The system cannot self-submit.

Cannot learn from biased data subsets
Pattern analysis requires minimum 50 samples per cluster-level combination before conclusions are applied. Outcomes from fewer than 50 applications are recorded but do not modify system behaviour.

SECTION 10: EXECUTION PROTOCOL — HOW TO RUN RESUME OS

This section is the operational guide. It defines exactly what to do, in what order, with what inputs, to produce a compliant Resume OS output.

10.1 Required Inputs

Master Resume	The candidate's complete, factual work history in any format. This is the source of truth. Nothing is added to it. Everything in the output must trace back to it.
Job Description or Job Link	The full JD text, OR a direct URL to the job posting. If URL: ATS Detection runs first. If JD text only: generic ATS-safe rules apply.
Optional — Candidate Preference	Tone preference (technical/formal/assertive/narrative), preferred variant structure, any roles or experiences to de-emphasise.

10.2 Execution Sequence — Step by Step

1. Step 1: Parse Job Input — Extract JD text (if URL, fetch and parse). Identify ATS from URL signature if present.
2. Step 2: Run Role Intelligence — Classify Role DNA: Function + Level + Environment + KPIs from JD language.
3. Step 3: Run ATS Intelligence — Apply ATS-specific keyword, structure, and formatting rules. Set file format output (PDF/DOCX).
4. Step 4: Run Recruiter Decision Model — Apply 6-second scan check. Flag all risk signals in master resume.
5. Step 5: Run Gap Analysis — Compare JD signals vs master resume signals. Produce gap scorecard.
6. Step 6: Run Optimization Engine — Apply all 8 sub-engines in sequence to working copy of master resume.
7. Step 7: Generate Variants — Produce 7-10 structural variants using the matched cluster template specifications.
8. Step 8: Score Variants — Calculate ATS Compatibility, Recruiter Scan, Signal Density, and Role Match scores per variant.
9. Step 9: Produce Strategy Log — Document every change, every gap, every recommendation.
10. Step 10: Deliver to Candidate — Present variants with scores, default recommendation, and strategy log.

10.3 Output Package — What Is Delivered

- 7-10 resume variants as downloadable files (PDF and/or DOCX based on ATS requirement)
- Score table: each variant's ATS score, recruiter score, signal density, role match score
- Default recommendation: the single best variant for this specific application with rationale
- Strategy log: structured record of every optimisation decision made
- Gap notification: list of keywords or skills not present in master resume, framed as candidate development signals
- ATS advisory: detected ATS platform, specific rules applied, file format recommendation

10.4 Things the System Will Always Do

- Preserve every factual claim from the master resume unchanged
- Apply level-appropriate language regardless of what the candidate originally wrote
- Eliminate all detected risk signals from every output variant
- Produce a strategy log documenting every change
- Score every variant before delivery
- Recommend the single best variant for the specific application
- Notify the candidate of any structural skill gaps that cannot be resolved through rephrasing

10.5 Things the System Will Never Do

- Add experience, skills, metrics, or achievements not in the master resume
- Claim a seniority level inconsistent with actual evidence
- Submit a resume on behalf of the candidate without explicit approval
- Use formatting that fails the ATS Parsing Gate — no exceptions
- Recommend a variant with a score below 70% ATS Compatibility without a prominent warning
- Apply learning patterns derived from fewer than 50 data points in a cluster-level combination

SECTION 11: GLOSSARY AND REFERENCE DEFINITIONS

Role DNA	The structured classification output of Layer 1. Format: Function + Level + Environment + KPIs. Every resume optimisation run begins with a Role DNA.
ATS (Applicant Tracking System)	Software used by companies to receive, sort, and rank job applications. Does not primarily auto-reject resumes — ranks and deprioritises them by relevance score. Used by 98.4% of Fortune 500 companies (Jobscan 2024).
Parsing Gate	The ATS formatting check. If the resume fails parsing (tables in body, multi-column layout, embedded images, non-standard headers), all keyword and relevance scoring is invalidated. Binary: Pass or Fail.
Gap Scorecard	The structured output of Layer 4 (Gap Analysis). Five dimensions: Keyword Coverage %, Title Alignment, Role Match Score, Leadership Signal strength, Risk Level.
Strategy Log	The required documentation output produced with every Resume OS run. Records every change made, the rationale, the gap findings, and the variant recommendation.
Master Resume	The candidate's complete, factual, unedited work history. This is the source of truth. It is never overwritten. All Resume OS outputs are generated from working copies.
Signal Density	A measure of how much verifiable proof — metrics, scale indicators, leadership language — appears per unit of resume content. High signal density = recruiters can verify capability without effort.
Recency Weighting	The Resume OS principle that recent experience receives more bullets, more detail, and more keyword density than older experience. Older roles compress progressively.
Environment Adaptation	The sub-engine that adjusts tone and language framing based on whether the target company is a startup (speed/ownership), SME (versatility), or MNC (process/scale).
Variant	One of 7-10 structurally distinct resume outputs generated per Resume OS run. All variants contain identical factual content. They differ in structure, section order, tone, and signal emphasis.
ATS Compatibility Score	The simulated ATS ranking score for a resume variant. Calculated using: Keyword Match 35% + Title Match 25% + Experience Relevance 20% + Skills 10% + Parsing Gate (Pass/Fail).
Recruiter Scan Score	The score measuring how well a resume passes the 6-second human scan check: Clarity + Relevance + Proof + Risk, evaluated in the top 40% of the resume.
Picks	PickCV platform designation for candidates scoring 85% or above on the Candidate Ranking Engine composite score for a specific role. Role-specific. Visually highlighted in recruiter interface.
Risk Signal	Any element of a resume that increases the recruiter's perceived risk of making a hire. Includes vague ownership, no metrics, inconsistent trajectory, over-generalisation, and 'responsible for' language.
Bullet Engine	Resume OS Sub-Engine 5.3. Applies the cluster-calibrated and level-calibrated bullet formula to every experience entry. The highest-impact optimisation engine for recruiter scoring.
SAR Formula	Situation-Action-Result. The mandatory bullet format for Cluster 9 (Operations/Consulting). Mirrors case interview evaluation logic.

RESUME OS — VERSION 1.0

PickCV Platform | Bhoomer Digital Solutions Pvt. Ltd.

"Talent is not missing. It is invisible to the systems that decide who gets hired."

Resume OS exists to fix that. One application at a time.

connect@pickcv.com | +91 86390 82763

March 2026